

TorBrix: Blocking-Resistant and Privacy-Preserving Tor Bridge Distribution

Abstract

Tor is currently the most popular network for anonymous Internet communication. It critically relies on volunteer nodes called *bridges* for relaying Internet traffic when a user's ISP blocks connections to Tor. Unfortunately, current methods for distributing bridges are vulnerable to malicious users who obtain and block bridge addresses. In this paper, we propose TorBrix, a protocol for privacy-preserving distribution of Tor bridges to n users, even when an unknown number $t < n$ of these users are controlled by a malicious adversary. TorBrix distributes $O(t \log n)$ bridges and guarantees that all honest users can connect to Tor with high probability after $O(\log t)$ rounds of communication with the distributor.

1 Introduction

Mass surveillance and censorship increasingly threaten democracy and freedom of speech. A growing number of governments around the world restrict access to the Internet to protect their domestic political, social, financial, and security interests [25, 22]. Countering this trend is the rise of anonymous communication systems, which strive to counteract censorship and preserve the privacy of individuals in cyberspace. Tor [11] is the most popular of such systems with more than 2.5 million users on average per day [2]. Tor relays Internet traffic via more than 6,500 volunteer nodes called *relays* spread across the world [3]. By routing data through random paths in the network, Tor can protect the private information of its users such as identity, geographical location, and content accessed.

Since the list of all relays is publicly available, state-sponsored organizations can force Internet service providers (ISPs) to block access to them, making Tor unavailable in territories ruled by the state. When access to Tor is blocked, users have the option to use *bridges*, which are volunteer relays not listed in Tor's public directory [10]. Bridges serve only as entry points into the rest of the Tor network, and their addresses are carefully distributed to the users, with the hope that they cannot all be learned by censors. As of March 2016, about 3,000 bridge nodes were running daily in the Tor network [5].

Currently, bridges are distributed to users based on different strategies such as CAPTCHA-enabled email-based distribution and IP-based distribution [10]. Unfortunately, censors are using sophisticated attacks to obtain and block bridges, rendering Tor unavailable for many users [9, 17, 27]. Also, state of the art techniques

for bridge distribution either (1) cannot provably guarantee that *all* honest users¹ can access Tor [26, 19, 24]; (2) can only work when the number of corrupt users is known in advance [18]; (3) require fully trusted distributors [19, 18, 24]; and/or (4) cannot resist malicious attacks from the distributors [26, 19, 18, 24].

It is often crucial to guarantee Tor access for *all* honest users, even after waiting for a specific (but predictable) amount of time. With the current bridge distribution mechanism, a user may need to wait for an indefinite amount of time before it can find a bridge that can connect it to the Tor network. Even if a large fraction of the bridges remain unblocked, this unreliable access mechanism can discourage the user from using Tor in the future.


To provide such a guaranteed access, one may choose to rely on *a priori* knowledge about the censor's budget, the amount of resources such as corrupt users that the censor owns. It is, however, often challenging in practice to identify or even estimate the number of corrupt users, due to the sophisticated and highly variable nature of Internet censorship in many countries such as China [1, 13]. Several reports show that censorship activities increase significantly during upheavals such as mass protests and elections [?, ?] probably due to the high costs of continuously discovering and blocking anti-censorship mechanisms.

Moreover, it is desirable for a system to be robust to attacks on the distributor servers. For example, powerful adversaries often supported by governments and large corporations can corrupt these systems, not only to break users' anonymity by obtaining bridge assignment information but also to prevent the bridge distribution protocol from achieving its goals.

In this paper, we propose TorBrix, a bridge distribution protocol that guarantees Tor access for all honest users with high probability, without requiring any *a priori* knowledge about the number of corrupt users. TorBrix guarantees that the number of rounds until all honest users can connect to Tor is bounded by $O(\log t)$, where $t < n$ is the number of corrupt users. This is achieved by distributing at most $O(t \log n)$ bridge addresses, where n is the total number of users.

Our protocol adaptively increases the number of bridges distributed among the users with respect to the number of bridges recently blocked by the adversary. This reduces the number of bridges used by the proto-

¹By honest users, we mean the users that are not controlled by the censor to obtain the bridge addresses assigned to them.



images/model.pdf

Figure 1: Our network model

col at the expense of a small (logarithmic) latency cost, which is also a function of the adversary’s cost. As a result, the overhead of TorBrix will always be proportional to the observed amount of corruption by the adversary.

In practice, users join and leave the system frequently. Adding new users to the system while offering provable robustness against an unknown number of corrupt users is challenging. This is because either (1) the adversary can cause a denial of service to the new users if the protocol’s “next move” is based on the adversary’s behavior, or (2) the protocol cannot guarantee every user receives a usable bridge. We propose a simple technique to handle this with small (constant) latency overhead. The details are described in Section 3.3.

Our technique for computing user-bridge assignments does not depend on actual bridge addresses. Thus, the distributor can assign “bridge pseudonyms” to the users. This prevents an honest-but-curious distributor from snooping on the user-bridge assignments. We also show how to distribute these pseudonyms among a group of geographically-dispersed servers who can collectively give the users the information needed to reconstruct their bridge addresses. This protects the anonymity of the users against colluding distributors. The details are described in Section ??.

We also describe a privacy-preserving bridge distribution mechanism for the scenarios where a certain fraction of the distributors may be controlled maliciously by the adversary.

We stress that TorBrix can run independently from Tor so that the Tor network can focus on its primary purpose of providing anonymity.

The rest of this paper is organized as follows. In Section 1.1, we describe our network and threat model. In Section 1.2, we state our main result as a theorem. We review related work in Section 2. In Section 3, we describe our protocol for reliable bridge distribution; we start from a basic protocol and improve it as we continue. We describe our implementation of TorBrix and our simulation results in Section 5. Finally, we summarize and state our open problems in Section 6.

1.1 Our Model

In this section, we first define a basic model, where a single distributor performs the bridge distribution task, and then define a multiple distributors model, where a group of distributors collectively run our protocol. Figure 1 depicts our high-level network model.

Basic Model. We assume there are n users (or *clients*) who need to obtain bridge addresses to access Tor. Initially, we assume a single trusted server called the *bridge distributor* (or simply the *distributor*), which has access to a reliable supply of bridge addresses.

We assume an adversary (or *censor*) who can view the internal state and control the actions of t of the clients; we call these adversarially-controlled clients *corrupt users*. The adversary is *adaptive* meaning that it can corrupt users at any point of the protocol, up to the point of taking over t users.

The corrupt users have the ability to *block* bridges whose IP addresses they receive. A bridge that is blocked cannot be used by any user. The adversary does not have to block a bridge as soon as it finds its address; he is allowed to strategically (perhaps by colluding with other corrupt users) decide when to block a bridge. We refer to the other $n - t$ users as *honest users*. Each honest user wants to obtain a list of bridge addresses, at least one of which is not blocked and hence can be used to connect to Tor. We further assume that the adversary has no knowledge of the private random bits used by our protocol.

We make the standard assumption that there exists a rate-limited channel such as email that the users can use to send their requests for bridges to the distributor, but which is not suitable for interactive Internet communication such as web surfing.² The distributor runs our bridge distribution protocol locally and sends bridge assignments back to the users via the same channel.

Bridge Reachability. We assume that the distributors obtain the availability information of bridges (i.e., which bridges are blocked and which bridges are not) from the Tor network. This assumption relies on a technique for testing reachability of bridges from outside the censored territory. Dingledine [8], Ensafi et al. [14], and Burnett

²Completely blocking a service such as email would usually impose significant economic and political consequences for censors.

and Feamster [?] describe active scanning mechanisms that can be used to test the reachability of bridges in real time with low latency, and we expect that one of these mechanisms to be deployed on the Tor network in the near future. The details of these methods and their current challenges are out of the scope of our paper.

1.2 Our Result

Below is our main theorem, which we prove in Section 3.

Theorem 1. *There exists a bridge distribution protocol that can run among m distributors and guarantee the following properties with probability $1 - 1/n^c$, for some constant $c \geq 1$, in the presence of a malicious adversary corrupting at most $\lfloor m/3 \rfloor$ of the distributors:*

1. *All honest users can connect to Tor after $\lceil \log \lceil (t+1)/32 \rceil \rceil + 1$ rounds of communication with the distributors;*
2. *The total number of bridges required is $O(t \log n)$;*
3. *Each user receives m messages in each round;*
4. *Each distributor sends/receives $O(m^2 + n)$ messages;*
5. *Each message has length $O(\log n)$ bits.*

We also simulated a proof-of-concept prototype of TorBrix to measure the running time and bridge cost of the protocol. We discuss our simulation results in Section 5.

2 Related Work

Proxy Distribution. The bridge distribution problem has been studied as the *proxy distribution*, where a set of proxy servers outside a censorship territory are distributed among a set of users inside the territory. These proxies are used to relay Internet traffic to blocked websites.

Feamster et al. [15] propose a proxy distribution protocol that requires every user to solve a cryptographic puzzle to discover a proxy. To protect against the censor discovering a large fraction of the proxies, each puzzle should be difficult enough (e.g., a day per puzzle). Unfortunately, this requires honest users to wait several hours and waste a large amount of computation power to solve the puzzles and receive a proxies even if there is no corruption in the network. Moreover, if the censor is computationally-powerful (as is the case for most governments and many companies), it can discover a large fraction of the proxies.

The Kaleidoscope system of Sovran et al. [24] disseminates proxy addresses over a social network whose links correspond to existing real world social relationships among users. Unfortunately, this protocol assumes the existence of a few internal trusted users who can relay

other users' traffic. Also, Kaleidoscope cannot guarantee its users' access to Tor.

McCoy et al. [19] propose Proximax; a proxy distribution system that uses social networks such as Facebook as trust networks that can provide a degree of protection against discovery by censors. Proximax estimates each user's effectiveness, and chooses the most effective users for advertising proxies, with the goals of maximizing the usage of these proxies while minimizing the risk of having them blocked.

Mahdian [18] studies the proxy distribution problem when the number of corrupt users, t , is known in advance. He proposes protocols for both large and small values of t and provides a lower bound for dynamic proxy distribution that is useful only when $t \ll n$. Unfortunately, it is usually hard in practice to reliably estimate the value of t . Mahdian's algorithm for large known t requires at most $t(1 + \lceil \log(n/t) \rceil)$ bridges, and his algorithm for small known t uses $O(t^2 \log n / \log \log n)$ bridges.

Wang et al. [26] propose a reputation-based bridge distribution mechanism called rBridge that computes every user's reputation based on the uptime of its assigned bridges and allows the user to replace a blocked bridge by paying some reputation credits. Interestingly, rBridge is the first model to provide user privacy against an honest-but-curious distributor. This is achieved by performing oblivious transfer between the distributor and the users along with commitments and zero-knowledge proofs for achieving unlinkability of transactions.

Handling DPI and Active Probing. The Tor Project has developed a variety of tools known as *pluggable transports* [4] to obfuscate the traffic transmitted between clients and bridges. This makes it hard for the censor to perform *deep packet inspection (DPI)* attacks, since distinguishing actual Tor traffic from legitimate-looking obfuscated traffic is hard.

The censor can also block bridges using *active probing*: he can passively monitor the network for suspicious traffic, and then actively probe dubious servers to block those that are determined to run the Tor protocol [13]. Depending on the sophistication of the censor, TorBrix may be used in parallel with tools that can handle DPI and active probing to provide further protection against blocking.

Resource-Competitive Analysis. Our analytical approach to bridge distribution can be seen as an application of the *resource-competitive analysis* introduced by Gilbert et al. [16, 6]. This approach evaluates the performance of any distributed protocol under attack by an adversary in the following way: if the adversary has a budget of t , then the worst-case resource cost of the algorithm is measured by some function of t . The adver-

sary’s budget is frequently expressed by the number of corrupt nodes controlled by the adversary. This model allows the system to adaptively increase/decrease its resource cost with the *current* amount of corruption by the adversary. Inspired by this model, we design resource-competitive algorithms for bridge distribution that scale reasonably with the adversary’s budget.

3 Our Protocol

We first construct a bridge distribution protocol that runs locally by a trusted distributor. Then, we extend this protocol to multiple distributors, where no distributor (nor any 1/3 coalition of them) learns any information about the user–bridge assignments.

We say an event occurs *with high probability*, if it occurs with probability at least $1 - 1/n^c$, for some constant $c \geq 1$. We denote the set of integers $\{1, \dots, n\}$ by $[n]$, the natural logarithm of any real number x by $\ln x$, and the logarithm to the base 2 of x by $\log x$. We denote a set of n users participating in our protocol by $\{u_1, \dots, u_n\}$. We define the *latency* of our protocol as the maximum number of rounds of communication that any user has to perform with the distributor(s) until he obtains at least one unblocked bridge.

3.1 Basic Protocol

The most naive approach to distributing a set of bridges to a set of users is to assign a unique bridge to each user. Tor relies on volunteer nodes that donate their bandwidth as relays and bridges to the Tor network. While the number of bridge users has nearly tripled in the past two years to about 50,000 [?], the number of bridges in the network has at best remained the same (about 3,000) over the same period of time [5]. Since the number of users is often much larger than the number of bridges, only a small group of users would receive bridges.

Thus, TorBrix assigns each bridge to multiple users. In particular, we start with a “small” set of bridges, and assign each user a bridge selected uniformly at random from this set. But how do we choose the size of this set? If the set is too small, an adversary can corrupt a small number of bridges and easily prevent any users from accessing Tor. If the set is too large, then we are wasting precious bridges.

The key idea is to *adjust* the number of bridges distributed in each round based on the number of bridges that have been blocked. Our protocol is divided into rounds incremented by i . We advance to the next round when the number of bridges blocked in the current round (b_i) exceeds a geometrically increasing threshold. In each round, we increase geometrically the size of the set of bridges that we assign (this size is the value d_i). In this way, we can ensure that the number of bridges TorBrix uses is a slowly growing function of the number

Algorithm 1 TorBrix – Basic Protocol

Goal: Distributes a set of $O(t \log n)$ bridges among a set of users $\{u_1, \dots, u_n\}$.

Run $3 \log n$ instances of the following algorithm in parallel with disjoint sets of bridges:

```

1:  $i \leftarrow 1$ 
2: while true do
3:    $d_i \leftarrow 2^{i+4}$ 
4:    $\{B_1, \dots, B_{d_i}\} \leftarrow d_i$  unblocked bridges
5:   for each  $j \in [n]$  do ▷ Distribute  $d_i$  bridges
6:     Pick  $k \in [d_i]$  uniformly at random
7:     Assign bridge  $B_k$  to user  $u_j$ 
8:   end for
9:   while  $b_i < 0.6 \times 2^{i+4}$  do
10:     $b_i \leftarrow \#$  blocked bridges in  $\{B_1, \dots, B_{d_i}\}$ 
11:   end while
12:    $i \leftarrow i + 1$ 
13: end while
```

of bridges blocked, which is a slowly growing function of the number of bad users t .

There are several technical issues remaining that must be addressed to ensure all users receive an unblocked bridge with high probability, and that we achieve the resource bounds of Theorem 1. Next, we describe our design in detail, and how we address those issues.

Our protocol (shown in Algorithm 1) is run locally by one distributor.³ TorBrix proceeds in *rounds* indicated by increments of the variable i in the while loop. In each round, the protocol recruits d_i bridges and assigns each user randomly to one of these bridges. Then, it continuously scans the set of bridges to count the number of blocked bridges, b_i . If this number exceeds a threshold, then the protocol proceeds to the next round.

The number of bridges distributed in every round is determined based on the threshold in that round as depicted in Figure 2. The exponential growth of the number of bridges distributed in each round allows us to achieve a logarithmic latency (on t) until all users can connect to Tor with high probability. In Lemma 3, we calculate the exact number of rounds required to achieve this goal. In Lemma 1, we show that if one instance of steps 1–13 of Algorithm 1 is executed, then it guarantees that all users can connect to Tor with some *constant probability*. Therefore, if we run $3 \log n$ instances in parallel, we can guarantee that all users connect to Tor *with high probability*.

In every round, TorBrix only distributes unblocked

³By “run locally”, we mean the distributor computes user-bridge assignments independent of any other distributor and without exchanging any information with them.

bridges. To reduce the total number of bridges required, we use unblocked bridges from previous rounds in the current distribution round. This can be done by removing blocked bridges from the pile of previously recruited bridges and adding a sufficient number of new bridges to accommodate the new load. One may choose to further reduce the number of bridges used by assigning new bridges only to those users whom one or more of their bridges have been blocked since the previous round. We skip this here to keep our algorithm and its analysis simple.

In each round, TorBrix sends to every user a single message containing $3 \log n$ bridge addresses assigned to this user in all instances of Algorithm 1 that run in parallel. This message is sent to the user via the rate-limited channel (e.g., email).

In the unlikely case that the censor blocks a significant number of the bridges such that the number of bridges to be distributed over all $3 \log n$ instances exceeds the number of users, n , then it becomes more reasonable to assign each user a unique bridge. This avoids distributing more than n bridges, which is overkill. Algorithm 1 can be modified to add an if-statement after Line 3 to check if $d_i \geq \frac{n}{3 \log n}$. If this is true, then the algorithm trivially assigns a unique bridge to every user and terminates. Otherwise, it executes lines 4–8. Note that this happens only if the adversary blocks a significant number of bridges, which we believe does not occur in most practical cases.

3.2 Handling Serialization Attacks

If the $3 \log n$ instances run completely independently, then the adversary can take advantage of this to increase the latency of the algorithm by a factor of $3 \log n$ using a *serialization attack*. In this attack, the adversary can strategically coordinate with its corrupt users to block the assigned bridges in such a way that the instances proceed to the next round one at a time. TorBrix prevents this attack by maintaining a single round counter, i , for all instances: whenever the number of blocked bridges in *any* of the instances exceeds the threshold for the current round, all instances are taken to the next round. Since all instances are run by the distributor locally, i can be easily synchronized between them. In Section ??, we describe how to modify Algorithm 1 to synchronize i .

3.3 Handling Client Churn

Algorithm 1 can only distribute bridges among a fixed set of users. A more realistic scenario is when users join or leave the algorithm frequently. One way to handle this is to add the new users to the algorithm from the next round (i.e., an increment of i). This, however, introduces two challenges. First, the number of corrupt users is unknown, and hence the adversary can arbitrarily delay the next round, causing a denial of service attack. Second,

our proof of robustness (Lemma 1) does not necessarily hold if n changes, because the algorithm is repeated $3 \log n$ times to ensure it succeeds with high probability.

To resolve these challenges, TorBrix assigns $3 \log n$ random bridges from the set of bridges recruited in the last round (i.e., the last time i was incremented) to every user who joins the protocol. This guarantees that new users are always assigned bridges once they join the system. If the total number of users, n , is doubled since the last round, recruit $3 \times 2^{i+4}$ unblocked bridges and assign 3 of them randomly to each user. This ensures that the number of parallel instances always remains $3 \log n$ even if n changes, because $\log n$ is increased by one when n is doubled. Therefore, each existing user must receive 3 new bridges so that Lemma 1 holds in the setting with churn. Our previous lemmas hold if users leave the system; thus, we only need to update n when nodes leave.

Finally, since distributing new bridges among existing users is done only after the number of users is doubled, the latency is increased by at most a $\log n$ term, where n is the largest number of users in the system during a complete run of the algorithm.

3.4 Privacy-Preserving Bridge Distribution

We now consider a *multiple distributors model*, where a group of $m \ll n$ distributors collectively distribute bridge addresses among the users such that none of the distributors can learn any information about the user–bridge assignments. Our goal is to run a protocol jointly among multiple distributors to keep user–bridge assignments hidden from each distributor and any coalition of up to a $1/3$ fraction of them. We assume that a sufficient number of bridges have already registered their addresses in the system.

We assume that the distributors are connected to each other pairwise via a synchronous network with reliable and authenticated channels. In this model, the adversary not only can corrupt an unknown number of the users, t but can also maliciously control and read the internal state of up to $\lfloor m/3 \rfloor$ of the distributors. The corrupt distributors can deviate from our protocols in any arbitrary manner, e.g., by sending invalid messages or remaining silent.

In each round i , the leader requests a group of at most d_i bridges to secret-share their IP addresses among all distributors (including the leader) using Shamir’s scheme [23].

Let (B_1, \dots, B_{d_i}) denote the sequence of shares the leader receives once the bridges finish the secret sharing protocol. The leader runs Algorithm 1 locally to assign B_j ’s to the users randomly, for all $j \in [d_i]$. Then, the leader broadcasts the pair (u_k, I_k) to all distributors, where I_k is the set of indices of bridges assigned to user u_k , for all $k \in [1, \dots, n]$.

Figure 2: Number of bridges distributed in round i of Algorithm 1. S and U indicate successful and unsuccessful iteration of the while-loop in Algorithm 1. An iteration is called successful when *all* users are able to connect to Tor in that iteration. Otherwise, it is called an unsuccessful iteration.

We first construct a *leader-based protocol*, where an honest-but-curious distributor called the *leader* locally runs Algorithm 1 over anonymous bridge addresses. The leader then sends anonymous user-bridge assignments to other distributors who can collectively “open” the assignments for the users.

Next, we construct a fully decentralized protocol, where a group of m distributors collectively compute the bridge distribution functionality while resisting malicious fault from up to a $\lfloor m/3 \rfloor$ fraction of the distributors. Malicious distributors not only may share information with other malicious entities but also can deviate from our protocol in any arbitrary manner, e.g., by sending invalid messages or remaining silent.

Both of these protocols rely on a secret sharing scheme for the bridges to share their IP addresses with the group of distributors. Before proceeding to our protocols, we briefly describe the secret sharing scheme used in our protocol.

Secret Sharing. A *secret sharing* protocol allows a party (called *the dealer*) to share a secret among m parties such that any set of τ or fewer parties cannot gain any information about the secret, but any set of at least $\tau + 1$ parties can reconstruct it. Shamir [23] proposed a secret sharing scheme, where the dealer shares a secret s among m parties by choosing a random polynomial $f(x)$ of degree τ such that $f(0) = s$. For all $j \in [m]$, the dealer sends $f(j)$ to the j -th party. Since at least $\tau + 1$ points are required to reconstruct $f(x)$, no coalition of τ or less parties can reconstruct s . The reconstruction algorithm requires a Reed-Solomon decoding algorithm [21] to correct up to $1/3$ invalid shares sent by dishonest distributors. In our protocols, we use the error correcting algorithm of Berlekamp and Welch [7].

4 Protocol Analysis

We now prove that Algorithm 1 achieves the properties described in Theorem 1 in the single distributor model. Although the adversary can corrupt up to t users, only some of the corrupt users might be actively blocking bridges in any given round. Before stating our first lemma, we define the following variables:

- b_i : number of bridges blocked in round i .
- d_i : number of bridges distributed in round i .
- t_i : number of corrupt users that have blocked at least one bridge in round i .

Lemma 1 (Robustness). *In round i of Algorithm 1, if $b_i < 0.6 \times 2^{i+4}$, then all honest users can connect to Tor with high probability.*

Proof. We first consider the execution of only one of the $3 \log n$ repeats of Algorithm 1. For each user, the algorithm chooses a bridge independently and uniformly at random and assigns it to the user. Without loss of generality, assume the corrupt users are assigned bridges first.

For $k = 1, 2, \dots, t_i$, let $\{X_k\}$ be a sequence of random variables each representing the bridge assigned to the k -th corrupt user. Also, let Y be a random variable corresponding to the number of *bad* bridges (i.e., the bridges that are assigned to at least one corrupt user) after all t_i corrupt users are assigned bridges. The sequence $\{Z_k = E[Y|X_1, \dots, X_k]\}$ defines a Doob martingale [12, Chapter 5], where $Z_0 = E[Y]$. Since each corrupt user is assigned a fixed bridge with probability $1/d_i$, the probability that the bridge is assigned to at least one corrupt user is $1 - (1 - 1/d_i)^{t_i}$. By symmetry, this probability is the same for all bridges. Thus, by linearity of expectation,

$$E[Y] = (1 - (1 - 1/d_i)^{t_i}) d_i < (1 - e^{-(t_i+1)/d_i}) d_i.$$

We know $t_i < 2^{i+4}$, because in each round $d_i = 2^{i+4}$ bridges are distributed and each corrupt user is assigned exactly one bridge; thus, each corrupt user can block at most one bridge. Hence,

$$E[Y] < (1 - 1/e^{1+1/2^{i+4}}) d_i \leq (1 - 1/e^2) d_i \quad (1)$$

Therefore, in expectation at most a constant fraction of the bridges become bad in each instance of the algorithm.

Since $|Z_{k+1} - Z_k| \leq 1$, $Z_0 = E[Y]$, and $Z_{t_i} = Y$, by the Azuma-Hoeffding inequality [12, Theorem 5.2],

$$\Pr(Y > E[Y] + \lambda) \leq e^{-2\lambda^2/t_i},$$

for any $\lambda > 0$. By setting $\lambda = \sqrt{d_i}$, we have

$$\Pr(Y > E[Y] + \sqrt{d_i}) \leq e^{-2d_i/t_i} < 1/e^2. \quad (2)$$

The last step holds since $t_i < d_i$. Therefore, with at most a constant probability, the actual number of bad bridges is larger than its expected value by at most $\sqrt{d_i}$. Therefore, the probability that an honest user is assigned a bad bridge is at most

$$\begin{aligned} \frac{E[Y] + \sqrt{d_i}}{d_i} &< \frac{(1 - 1/e^{1+1/2^{i+4}}) d_i + \sqrt{d_i}}{d_i} \\ &= 1/e^{1+1/2^{i+4}} + 1/\sqrt{d_i}, \end{aligned} \quad (3)$$

where the first step is achieved using (1).

Now, let $p_1 = \Pr(Y > \mathbb{E}[Y] + \sqrt{d_i})$, and let p_2 be the probability that a fixed honest user is assigned a bad bridge in a fixed instance and a fixed round. From (2) and (3), we have

$$p_1 < 1/e^2 \quad \text{and} \quad p_2 < 1/e^{1+1/2^{i+4}} + 1/\sqrt{d_i}.$$

Thus, the probability that a fixed user fails to receive a good bridge in a fixed instance and a fixed round is equal to $p_1 + (1 - p_1)p_2$, which is at most 0.6.

If the algorithm is repeated $\lceil 3 \log n \rceil$ times in parallel, then the probability that the user is assigned to only bad bridges in the last round is at most $0.6^{\lceil 3 \log n \rceil} \leq 1/n^2$. By a union bound, the probability that any of the n users is assigned only bad bridges in a round is at most $1/n$. Therefore, all honest users can connect to Tor with high probability. \square

Algorithm 1 does not necessarily assign the same number of users to each bridge. However, in the following lemma, we show that each bridge is assigned to almost the same number of users as other bridges with high probability providing a reasonable level of load-balancing.

Lemma 2 (Bridge Load-Balancing). *Let X be a random variable representing the maximum number of users assigned to any bridge, Y be a random variable representing the minimum number of users assigned to any bridge, and $z = \Theta(\frac{\ln n}{\ln \ln n})$. Then, we have*

$$\Pr(X \geq \mu z) \leq 2/n \quad \text{and} \quad \Pr(Y \leq \mu z) \leq 2/n,$$

where $\mu = n/d_i$.

Proof. Each round of Algorithm 1 can be seen as the classic balls-and-bins process: n balls (users) are thrown independently and uniformly at random into d_i bins (bridges). It is well known that the distribution of the number of users assigned to a bridge is approximately Poisson with $\mu = n/d_i$ [20, Chapter 5].

Let X_j be the random variable corresponding to the number of users assigned to the j -th bridge, and let \tilde{X}_j be the Poisson random variable approximating X_j . We have $\mu = \mathbb{E}[X_j] = \mathbb{E}[\tilde{X}_j] = n/d_i$. We use the following Chernoff bounds from [20, Chapter 5] for Poisson random variables:

$$\Pr(\tilde{X}_j \geq x) \leq e^{-\mu} (e\mu/x)^x, \text{ when } x > \mu \quad (4)$$

$$\Pr(\tilde{X}_j \leq x) \leq e^{-\mu} (e\mu/x)^x, \text{ when } x < \mu \quad (5)$$

Let $x = \mu y$, where $y = ez$. From (4), we have

$$\begin{aligned} \Pr(\tilde{X}_j \geq \mu y) &\leq \left(\frac{e^{y-1}}{y^y} \right)^\mu \\ &\leq \frac{e^{y-1}}{y^y} \\ &= \frac{1}{e} \left(\frac{1}{z^z} \right)^e < \frac{1}{n^2}. \end{aligned} \quad (6)$$

The second step is because $y^y > e^{y-1}$ (since $z > 1$) and $\mu > 1$. The last step is because $z = \Theta(\frac{\ln n}{\ln \ln n})$ is the solution of $z^z = n$. To show this, we take log of both sides of $z^z = n$ twice, which yields

$$\ln z + \ln \ln z = \ln \ln n.$$

We have

$$\ln z \leq \ln z + \ln \ln z = \ln \ln n < 2 \ln z.$$

Since $z \ln z = \ln n$,

$$z/2 < \frac{\ln n}{\ln \ln n} \leq z.$$

Therefore, $z = \Theta(\frac{\ln n}{\ln \ln n})$.

It is shown in [20, Corollary 5.11] that for any event that is monotone in the number of balls, if the event occurs with probability at most p in the Poisson approximation, then it occurs with probability at most $2p$ in the exact case. Since the maximum and minimum bridge loads are both monotonically increasing in the number of users, from (6) we have

$$\Pr(X_j \geq \mu y) \leq 2\Pr(\tilde{X}_j \geq \mu y) < 2/n^2.$$

By applying a union bound over all bridges, the probability that the number of users assigned to any bridge will be more than μz is at most $2/n$. The bound on the minimum load can be shown using inequality (5) in a similar way. \square

Lemma 3 (Latency). *By running Algorithm 1, all honest users can connect to Tor with high probability after at most $\lceil \log \lceil (t+1)/32 \rceil \rceil + 1$ iterations of the while loop.*

Proof. Let k denote the smallest number of rounds required until all users can connect to Tor with high probability. Intuitively, k is bounded, because the number of corrupt nodes, t , is bounded. In the following, we find k with respect to t .

Without loss of generality, we only consider one of the $3 \log n$ parallel instances of Steps 1–13 of Algorithm 1. The best strategy for the adversary is to maximize k , because this prevents the algorithm from succeeding soon.⁴

⁴Because otherwise both the number of bridges used and the latency become smaller. As soon as the algorithm stops going to next round, it guarantees based on Lemma 1 that all honest users can connect to Tor with high probability.

In each round i , this can be achieved by minimizing the number of bridges blocked (i.e., b_i), while ensuring the algorithm proceeds to the next round. However, the adversary has to block at least $0.6 \times 2^{i+4}$ bridges in each round to force the algorithm to proceed to the next round. Let ℓ be the smallest integer such that $2^\ell \geq t$. In round ℓ , the adversary has enough corrupt users to take the algorithm to round $\ell + 1$. However, in round $\ell + 1$, the adversary can block at most $2^\ell < 2^{\ell+1}$ bridges, which is insufficient for proceeding to round $\ell + 2$. Therefore, $\ell + 1$ is the last round and $k = \ell + 1$. Since $2^\ell \geq t$, and the algorithm starts by distributing 32 bridges,

$$k = \lceil \log \lceil (t+1)/32 \rceil \rceil + 1.$$

In other words, if the while loop runs for at least $\lceil \log \lceil (t+1)/32 \rceil \rceil + 1$ iterations, then with high probability all honest users can connect to Tor. \square

Lemma 4 (Bridge Cost). *The total number of bridges used by Algorithm 1 is at most $\min[(10t + 96) \log n, 2n]$.*

Proof. Consider one of the $3 \log n$ instances of Algorithm 1. The algorithm starts by distributing 32 bridges. In every round $i > 0$, the algorithm distributes a new bridge only to replace a bridge blocked in round $i - 1$. Thus, the total number of bridges used until round i , denoted by M_i , is equal to the number of bridges blocked until round i plus the number of new bridges distributed in round i , which we denote by a_i . Therefore,

$$M_i = a_i + \sum_{j=0}^{i-1} b_j. \quad (7)$$

In round i , the algorithm recruits $a_i \leq 2^{i+4}$ new bridges, because some of the bridges required for this round might be reused from previous rounds. Since in round i we have $b_i < 0.6 \times 2^{i+4}$,

$$M_i < 2^{i+4} + 0.6 \sum_{j=1}^{i-1} 2^{j+4} = 9.6(2^i - 2) + 2^{i+4}$$

From Lemma 3, it is sufficient to run the algorithm $k = \lceil \log \lceil (t+1)/32 \rceil \rceil + 1$ rounds. Therefore,

$$M_k < 9.6(2^k - 2) + 2^{k+4} \leq 3.2t + 32.$$

Since the algorithm is repeated $3 \log n$ times, $M_k < (10t + 96) \log n$. When the number of bridges to be distributed in the current round across all instances becomes larger than n , the algorithm distributes exactly n bridges among the users. Therefore, the total number of bridges used is at most $\min[(10t + 96) \log n, 2n]$. \square

5 Evaluation

We implemented a proof-of-concept prototype of TorBrix and tested it in a simulated environment under various adversarial behavior. We set the parameters of TorBrix in such a way that we ensure it fails with probability at most 10^{-4} . We consider three blocking strategies for the adversary: *prudent*, *aggressive*, and *stochastic*. A prudent adversary blocks the minimum number of bridges in each round such that the algorithm is forced to go to the next round. An aggressive adversary blocks immediately all of the bridges he learns from the corrupt users. Finally, a stochastic adversary blocks each bridge he receives with some fixed probability.

To evaluate the performance of TorBrix, we calculate five measures of performance in two experiments. These measures are:

1. **Failed users:** Users who do not currently have any unblocked bridges.
2. **Bridges distributed:** Bridges distributed in the current round (d_i).
3. **Bridges blocked:** Bridges blocked in the current round (b_i).
4. **Bridges used:** The set of all unique bridges distributed by the algorithm until this round (M_i).
5. **Latency:** Number of rounds until all users receive at least one unblocked bridge.

In the first experiment (shown in Figure 3), we run the basic algorithm for $n = 65536$ and $t = 180$, and calculate Measures 1–4 at the end of each round after running the algorithm over ten samples for a fixed set of parameters. The experiment was run with the three different adversarial strategies. For the stochastic blocking, the adversary blocks each bridge with probability 0.95. In the second experiment (shown in Figure 4), we run the algorithm using a single distributor for $n = 1024$ and calculate measures 4 and 5 by varying t between 0 and 1023.

Our results indicate that TorBrix incurs a small cost when there is small or no corruption in the network. Moreover, the algorithm scales well with the number of corruptions and can quickly adapt to the adversary’s behavior. These all support our claim that TorBrix can be used for practical bridge distribution that guarantees access for all users.

6 Conclusion

We described TorBrix, a bridge distribution system that allows all honest users to connect to Tor in the presence of an adversary corrupting an unknown number of users. Our algorithm can adaptively increase the number of bridges according to the behavior of the adversary and it




Figure 3: Simulation results for $n = 65,536$ and $t = 180$ with prudent (left), aggressive (middle), and stochastic (right) adversary.

Figure 4: Simulation results for $n = 1024$ and variable number of corrupt users with prudent (left), aggressive (middle), and stochastic (right) adversary.

uses a near-optimal number of bridges. We also modified our algorithm slightly to handle user churn (join/leave) by adding small (constant) amortized latency.

Next, we described a protocol for privacy-preserving bridge distribution that runs our centralized distribution algorithm obviously among a group of distributors. We showed that the resulting protocol not only can protect the privacy of user-bridge assignments from any coalition of up to a $1/3$ fraction of the distributors but also can tolerate malicious attacks from a $1/3$ fraction of the distributors.

Although TorBrix represents a step towards robust and privacy-preserving bridge distribution, many challenges remain for future work. For example, the current algorithm uses a relatively large number of bridges when the number of corrupt users is large. Is it possible to make the bridge cost sublinear in t with practical constant terms?

An interesting direction is to use inexpensive honeypot bridges for detecting and blacklisting corrupt users. This, however, requires a mechanism such as CAPTCHA for preventing the adversary from distinguishing real bridges from the fake ones. Moreover, a colluding adversary may be able to compare bridges assigned to its corrupt users to detect honeypots.

To better explore the possibility of achieving a sublinear bridge cost, one may consider finding lower bounds for different scenarios. For example, when each user is assigned at least one bridge, it seems impossible to achieve a sublinear bridge cost unless some of the bridges are fake, or we only distribute real bridges in random-chosen rounds. What is the lower bound for the number of rounds in these scenarios? Another interest-

ing open problem is to examine if our current notion of robustness is overkill for practice. For example, is it possible to significantly reduce our costs by guaranteeing access for all but a constant number of users?

References

- [1] The Open Net Initiative: China, 2012.
- [2] The Tor Project metrics: Direct users connecting between January 1, 2015 and March 31, 2015, 2015.
- [3] The Tor Project metrics: Relays in the network between January 1, 2015 and March 31, 2015, 2015.
- [4] The Tor Project: Pluggable transport, 2015.
- [5] The Tor Project metrics: Bridges in the network between March 1, 2016 and March 31, 2016, 2016.
- [6] M. A. Bender, J. T. Fineman, M. Movahedi, J. Saia, V. Dani, S. Gilbert, S. Pettie, and M. Young. Resource-competitive algorithms. *ACM SIGACT News*, 46(3):57–71, Sept. 2015.
- [7] E. Berlekamp and L. Welch. Error correction for algebraic block codes, US Patent 4,633,470, Dec. 1986.
- [8] R. Dingledine. Research problem: Five ways to test bridge reachability, 2011.
- [9] R. Dingledine. Research problems: Ten ways to discover Tor bridges, 2011.

- [10] R. Dingledine and N. Mathewson. Design of a blocking-resistant anonymity system. Technical report, The Tor Project Inc., 2006.
- [11] R. Dingledine, N. Mathewson, and P. Syverson. Tor: the second-generation onion router. In *Proceedings of the 13th USENIX Security Symposium*, Berkeley, CA, USA, 2004.
- [12] D. P. Dubhashi and A. Panconesi. *Concentration of Measure for the Analysis of Randomized Algorithms*. Cambridge University Press, New York, NY, USA, 2009.
- [13] R. Ensafi, D. Fifield, P. Winter, N. Feamster, N. Weaver, and V. Paxson. Examining how the Great Firewall discovers hidden circumvention servers. In *Internet Measurement Conference (IMC)*. ACM, 2015.
- [14] R. Ensafi, J. Knockel, G. Alexander, and J. R. Crandall. Detecting intentional packet drops on the Internet via TCP/IP side channels. In *Proceedings of the 15th International Conference on Passive and Active Measurement - Volume 8362*, PAM 2014, pages 109–118, New York, NY, USA, 2014. Springer-Verlag New York, Inc.
- [15] N. Feamster, M. Balazinska, W. Wang, H. Balakrishnan, and D. Karger. Thwarting web censorship with untrusted messenger discovery. In R. Dingledine, editor, *Privacy Enhancing Technologies*, volume 2760 of *Lecture Notes in Computer Science*, pages 125–140. Springer Berlin Heidelberg, 2003.
- [16] S. Gilbert, J. Saia, V. King, and M. Young. Resource-competitive analysis: A new perspective on attack-resistant distributed computing. In *Proceedings of the 8th International Workshop on Foundations of Mobile Computing, FOMC '12*, pages 1:1–1:6, New York, NY, USA, 2012. ACM.
- [17] Z. Ling, J. Luo, W. Yu, M. Yang, and X. Fu. Extensive analysis and large-scale empirical evaluation of tor bridge discovery. In *INFOCOM, 2012 Proceedings IEEE*, pages 2381–2389, March 2012.
- [18] M. Mahdian. Fighting censorship with algorithms. In P. Boldi and L. Gargano, editors, *Fun with Algorithms*, volume 6099 of *Lecture Notes in Computer Science*, pages 296–306. Springer Berlin Heidelberg, 2010.
- [19] D. McCoy, J. A. Morales, and K. Levchenko. Proximax: Measurement-driven proxy dissemination. In *Proceedings of the 15th International Conference on Financial Cryptography and Data Security, FC'11*, pages 260–267, Berlin, Heidelberg, 2012. Springer-Verlag.
- [20] M. Mitzenmacher and E. Upfal. *Probability and Computing: Randomized Algorithms and Probabilistic Analysis*. Cambridge University Press, 2005.
- [21] I. Reed and G. Solomon. Polynomial codes over certain finite fields. *Journal of the Society for Industrial and Applied Mathematics (SIAM)*, pages 300–304, 1960.
- [22] D. Rushe. Google reports 'alarming' rise in censorship by governments. The Guardian, June 2012.
- [23] A. Shamir. How to share a secret. *Commun. ACM*, 22(11):612–613, 1979.
- [24] Y. Sovran, A. Libonati, and J. Li. Pass it on: Social networks stymie censors. In *Proceedings of the 7th International Conference on Peer-to-peer Systems, IPTPS'08*, pages 3–3, Berkeley, CA, USA, 2008. USENIX Association.
- [25] K. Turner. Mass surveillance silences minority opinions, according to study. The Washington Post, March 2016.
- [26] Q. Wang, Z. Lin, N. Borisov, and N. Hopper. rbridge: User reputation based tor bridge distribution with privacy preservation. In *Network and Distributed System Security Symposium, NDSS 2013*. The Internet Society, 2013.
- [27] P. Winter and S. Lindskog. How the great firewall of China is blocking Tor. In *2nd USENIX Workshop on Free and Open Communications on the Internet*, Berkeley, CA, 2012.